

ResearchMind

Learning Session Export

Topic: <https://arxiv.org/pdf/2409.09825v4> · Exported: 11 June 2026, 00:45 · Session ID: 035a029f

Overall Progress: 0%

Knowledge Profile

Level: —

Reading Sequence

1. Attention Is All You Need

Vaswani, Shazeer, Parmar et al. · 2017 · 87,000 citations · [Foundational]

We propose the Transformer, a model architecture based solely on attention mechanisms.

2. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

Devlin, Chang, Lee, Toutanova · 2018 · 52,000 citations · [Essential]

We introduce BERT, which obtains new state-of-the-art results on eleven NLP tasks.

3. Language Models are Few-Shot Learners

Brown et al. · 2020 · 38,000 citations · [Milestone]

We train GPT-3, an autoregressive language model with 175 billion parameters.

4. An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale

Dosovitskiy et al. · 2020 · 25,000 citations · [Extension]

We apply a pure transformer to sequences of image patches for image classification.

Spaced Repetition Schedule

Concept	Due In	Memory Strength
Scaled Dot-Product Attention	1 day	32%
Positional Encoding	3 days	55%
Multi-Head Attention	1 day	28%

Feed-Forward Sublayer	6 days	72%
Encoder–Decoder Stack	2 days	44%